

Assignment # 4

By

Rukhsar Jabeen (BSCS-8961)



Submitted to: Maam Sadaf Tanvir

Subject: STIC

Date: 12. dec.2025

Batch: Spring-2024

Semester: Fourth

DEPARTMENT OF COMPUTER SCIENCE

**ABASYN UNIVERSITY
ISLAMABAD CAMPUS**

Topic:

Hyperparameter Tuning in Adam Optimizer for MobileNetV2

◆ Abstract

This study examines how changes in Adam optimizer hyperparameters affect the performance of a neural network model. Five experiments were conducted, including a baseline configuration and four modified versions that adjusted β_1 , β_2 , ϵ , and label-smoothing values. All models were trained under the same conditions, and their training accuracy, validation accuracy, and total training time were recorded. The results show that altering β_1 , β_2 , and smoothing led to small improvements in validation accuracy, while increasing ϵ reduced it. Training accuracy remained high across all experiments, suggesting that the main differences were in generalization rather than in the model's ability to fit the training data. These findings illustrate the influence of optimizer settings on model stability and validation performance.

1. Introduction:

The **purpose** of this experiment is to examine how the Adam optimizer's hyperparameters β_1 , β_2 , ϵ , and the smoothing factor affect the performance of a MobileNetV2 model trained on the Real Field dataset.

- These parameters directly influence how gradients are updated during training, which can change both the speed of convergence and the overall accuracy of the model. Since Adam is commonly used with its default values, it is important to understand how modifying these settings can improve or weaken generalization.
- To investigate this, several experiments were carried out using a baseline configuration and multiple variations in which each hyperparameter was adjusted independently. By comparing training accuracy, validation accuracy, and total training time across these runs, the study **aims** to identify which hyperparameter changes have a meaningful impact on the model's behavior. This helps clarify the relationship between Adam's internal settings and the stability and effectiveness of the training process.

2. Mathematical Background:

- The Adam optimizer is a first-order gradient-based optimization method that combines ideas from Momentum and **RMSProp** to adapt learning rates for each parameter.

- It maintains estimates of both the first moment (the mean of the gradients) and the second moment (the uncentered variance), and uses these to update model parameters in a stable and efficient way. The main update steps are given below.

1. First moment update

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

2. Second moment update

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

3. Bias-corrected estimates

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

4. Parameter update rule

$$\theta_{t+1} = \theta_t - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

- Here, g_t is the gradient of the loss with respect to the parameter at time step t
- α is the learning rate
- β_1 and β_2 control how quickly the first and second moments decay.
- and ϵ is a small constant to ensure numerical stability. These components work together to provide smooth and adaptive updates, which often leads to faster convergence compared to standard gradient descent.

3. Experimental Setup

→ **Dataset:**

The experiments were conducted using the Real Field Dataset, with images organized in the following directories:

- Training: drive/MyDrive/dataset/train
- Validation: drive/MyDrive/dataset/Validation
- Test (optional): drive/MyDrive/dataset/test

```

# --- Custom Data Paths ---
TRAIN_PATH = "drive/MyDrive/dataset/train"
VAL_PATH = "drive/MyDrive/dataset/Validation"
TEST_PATH = "drive/MyDrive/dataset/test"

```

- The training set contains 6152 images and the validation set contains 596 images across 6 classes. Images were resized to 32×32 pixels and normalized to the $[0,1]$ range.
- Model: MobileNetV2 with Global Average Pooling + Dense layer with Num classes Outputs. Optimizer: Adam with fixed learning rate = 0.001.

→ Hyperparameter Settings:

The Adam optimizer was tested with five different configurations:

Baseline	$\beta_1=0.9$	$\beta_2=0.999,$	$\epsilon=1e-8$
High β_1	$\beta_1=0.95$	$\beta_2=0.999$	$\epsilon=1e-8$
Low β_2	$\beta_1=0.9$	$\beta_2=0.99$	$\epsilon=1e-8$
Large ϵ	$\beta_1=0.9$	$\beta_2=0.999$	$\epsilon=1e-6$
Very High Smoothing	$\beta_1=0.99$	$\beta_2=0.9999$	$\epsilon=1e-7$

- Training Configuration:**

All models were trained for 10 epochs with a fixed learning rate of 0.001 and batch size of 64. Training and validation accuracy were recorded at each epoch, along with total training time, to evaluate performance and efficiency.

4. Results and Analysis

▪ Baseline:

```

=====
Running experiment: baseline
Epoch 1/10
97/97      980s 10s/step - accuracy: 0.4375 - loss: 1.4625 - val_accuracy: 0.1275 - val_loss: 1.7910
Epoch 2/10
97/97      230s 2s/step - accuracy: 0.7283 - loss: 0.7135 - val_accuracy: 0.1275 - val_loss: 1.8022
Epoch 3/10
97/97      258s 2s/step - accuracy: 0.8275 - loss: 0.4605 - val_accuracy: 0.1275 - val_loss: 1.8603
Epoch 4/10
97/97      238s 2s/step - accuracy: 0.8779 - loss: 0.3310 - val_accuracy: 0.1275 - val_loss: 1.9432
Epoch 5/10
97/97      255s 2s/step - accuracy: 0.8836 - loss: 0.3283 - val_accuracy: 0.1275 - val_loss: 2.0603
Epoch 6/10
97/97      273s 3s/step - accuracy: 0.9328 - loss: 0.1965 - val_accuracy: 0.1275 - val_loss: 2.1921
Epoch 7/10
97/97      251s 2s/step - accuracy: 0.9043 - loss: 0.2699 - val_accuracy: 0.1275 - val_loss: 2.3730
Epoch 8/10
97/97      241s 2s/step - accuracy: 0.9434 - loss: 0.1575 - val_accuracy: 0.1275 - val_loss: 2.4309
Epoch 9/10
97/97      254s 2s/step - accuracy: 0.9645 - loss: 0.1173 - val_accuracy: 0.1275 - val_loss: 2.7153
Epoch 10/10
97/97      234s 2s/step - accuracy: 0.9648 - loss: 0.1041 - val_accuracy: 0.1275 - val_loss: 2.8570
Finished: baseline

```

- Training loss dropped steadily and accuracy reached **96.48%**, but validation accuracy stayed very low at **12.75%**, indicating overfitting.
- Validation loss increased, showing poor generalization. Training was slow (\approx 3215 s). The optimizer reduced training loss effectively, but the model failed to perform well on unseen data.

High_beta 1:

```
=====
Running experiment: high_beta1
Epoch 1/10
97/97 279s 2s/step - accuracy: 0.3888 - loss: 1.5351 - val_accuracy: 0.2366 - val_loss: 1.7889
Epoch 2/10
97/97 231s 2s/step - accuracy: 0.7270 - loss: 0.7631 - val_accuracy: 0.2366 - val_loss: 1.7816
Epoch 3/10
97/97 271s 2s/step - accuracy: 0.8267 - loss: 0.4836 - val_accuracy: 0.1275 - val_loss: 1.8387
Epoch 4/10
97/97 234s 2s/step - accuracy: 0.8892 - loss: 0.2898 - val_accuracy: 0.1275 - val_loss: 1.8743
Epoch 5/10
97/97 233s 2s/step - accuracy: 0.9093 - loss: 0.2355 - val_accuracy: 0.1275 - val_loss: 2.0187
Epoch 6/10
97/97 241s 2s/step - accuracy: 0.9147 - loss: 0.2289 - val_accuracy: 0.1275 - val_loss: 2.2058
Epoch 7/10
97/97 233s 2s/step - accuracy: 0.9243 - loss: 0.2254 - val_accuracy: 0.1275 - val_loss: 2.3894
Epoch 8/10
97/97 239s 2s/step - accuracy: 0.9279 - loss: 0.2052 - val_accuracy: 0.1275 - val_loss: 2.3120
Epoch 9/10
97/97 258s 2s/step - accuracy: 0.9581 - loss: 0.1169 - val_accuracy: 0.1275 - val_loss: 2.8143
Epoch 10/10
97/97 231s 2s/step - accuracy: 0.9759 - loss: 0.0728 - val_accuracy: 0.1275 - val_loss: 2.8871
Finished: high_beta1
```

- **Training accuracy** improved faster, reaching **97.59%**, and loss decreased steadily.
- **Validation accuracy** slightly improved initially (**23.66%**) but mostly stayed low (**12.75%**), showing overfitting similar to baseline. Training was faster (**~2452 s**) than baseline. High β_1 helped convergence speed but did not improve generalization.

Low_beta2:

```
=====
Running experiment: low_beta2
Epoch 1/10
97/97 279s 2s/step - accuracy: 0.4005 - loss: 1.5129 - val_accuracy: 0.2366 - val_loss: 1.7951
Epoch 2/10
97/97 232s 2s/step - accuracy: 0.6994 - loss: 0.7875 - val_accuracy: 0.2366 - val_loss: 1.8265
Epoch 3/10
97/97 239s 2s/step - accuracy: 0.7896 - loss: 0.5692 - val_accuracy: 0.1275 - val_loss: 1.8508
Epoch 4/10
97/97 255s 2s/step - accuracy: 0.8594 - loss: 0.3774 - val_accuracy: 0.1275 - val_loss: 1.9155
Epoch 5/10
97/97 234s 2s/step - accuracy: 0.8781 - loss: 0.3316 - val_accuracy: 0.1174 - val_loss: 2.0168
Epoch 6/10
97/97 240s 2s/step - accuracy: 0.9083 - loss: 0.2592 - val_accuracy: 0.1174 - val_loss: 2.0619
Epoch 7/10
97/97 254s 2s/step - accuracy: 0.9206 - loss: 0.2246 - val_accuracy: 0.1174 - val_loss: 2.1183
Epoch 8/10
97/97 241s 2s/step - accuracy: 0.9216 - loss: 0.2346 - val_accuracy: 0.1174 - val_loss: 2.0645
Epoch 9/10
97/97 235s 2s/step - accuracy: 0.9507 - loss: 0.1401 - val_accuracy: 0.1174 - val_loss: 2.1294
Epoch 10/10
97/97 234s 2s/step - accuracy: 0.9488 - loss: 0.1558 - val_accuracy: 0.1174 - val_loss: 2.1865
Finished: low_beta2
```

Training accuracy reached **94.88%** with steadily decreasing loss, but **validation accuracy** remained low (**~11.74%**), indicating overfitting. Training time was moderate (**~2443 s**). Lowering β_2 slightly sped up convergence but did not improve generalization.

large_epsilon:

```
=====
Running experiment: large_epsilon
Epoch 1/10
97/97 289s 3s/step - accuracy: 0.4162 - loss: 1.5047 - val_accuracy: 0.1174 - val_loss: 1.7900
Epoch 2/10
97/97 237s 2s/step - accuracy: 0.7166 - loss: 0.7537 - val_accuracy: 0.1174 - val_loss: 1.8045
Epoch 3/10
97/97 241s 2s/step - accuracy: 0.8064 - loss: 0.5054 - val_accuracy: 0.1174 - val_loss: 1.8589
Epoch 4/10
97/97 257s 2s/step - accuracy: 0.8456 - loss: 0.4178 - val_accuracy: 0.1174 - val_loss: 1.9146
Epoch 5/10
97/97 239s 2s/step - accuracy: 0.8787 - loss: 0.3329 - val_accuracy: 0.1174 - val_loss: 2.0072
Epoch 6/10
97/97 253s 3s/step - accuracy: 0.9132 - loss: 0.2544 - val_accuracy: 0.1174 - val_loss: 2.1124
Epoch 7/10
97/97 245s 3s/step - accuracy: 0.9378 - loss: 0.1751 - val_accuracy: 0.1174 - val_loss: 2.2285
Epoch 8/10
97/97 237s 2s/step - accuracy: 0.9533 - loss: 0.1407 - val_accuracy: 0.1174 - val_loss: 2.3416
Epoch 9/10
97/97 231s 2s/step - accuracy: 0.9601 - loss: 0.1081 - val_accuracy: 0.1174 - val_loss: 2.4208
Epoch 10/10
97/97 232s 2s/step - accuracy: 0.9650 - loss: 0.0988 - val_accuracy: 0.1174 - val_loss: 2.4715
Finished: large_epsilon
```

Training accuracy reached **96.50%** with smooth loss reduction, but **validation accuracy** remained very low (**~11.74%**), showing poor generalization. Increasing ϵ sped up training slightly but did not improve validation performance.

Very_high_smoothing:

```
=====
Running experiment: very_high_smoothing
Epoch 1/10
97/97 279s 2s/step - accuracy: 0.4017 - loss: 1.5167 - val_accuracy: 0.2366 - val_loss: 1.7883
Epoch 2/10
97/97 229s 2s/step - accuracy: 0.6547 - loss: 0.8882 - val_accuracy: 0.2366 - val_loss: 1.7848
Epoch 3/10
97/97 269s 2s/step - accuracy: 0.7476 - loss: 0.6742 - val_accuracy: 0.2366 - val_loss: 1.8208
Epoch 4/10
97/97 262s 2s/step - accuracy: 0.8262 - loss: 0.4634 - val_accuracy: 0.2366 - val_loss: 1.8474
Epoch 5/10
97/97 232s 2s/step - accuracy: 0.8951 - loss: 0.2744 - val_accuracy: 0.2366 - val_loss: 1.9332
Epoch 6/10
97/97 228s 2s/step - accuracy: 0.9249 - loss: 0.2247 - val_accuracy: 0.2366 - val_loss: 1.9823
Epoch 7/10
97/97 229s 2s/step - accuracy: 0.9210 - loss: 0.2227 - val_accuracy: 0.2131 - val_loss: 2.1686
Epoch 8/10
97/97 259s 2s/step - accuracy: 0.9343 - loss: 0.1974 - val_accuracy: 0.2131 - val_loss: 2.2453
Epoch 9/10
97/97 267s 2s/step - accuracy: 0.9398 - loss: 0.1773 - val_accuracy: 0.2131 - val_loss: 2.3435
Epoch 10/10
97/97 229s 2s/step - accuracy: 0.9453 - loss: 0.1487 - val_accuracy: 0.2131 - val_loss: 2.4959
Finished: very_high_smoothing
```

Training accuracy steadily improved to **94.53%**, with gradual loss reduction. **Validation accuracy** peaked at **23.66%** but slightly dropped later, indicating minor overfitting. This configuration converged slightly slower initially but maintained stable training.

4.1 Training and Validation Accuracy:

In short:

- **Baseline ($\beta_1=0.9$, $\beta_2=0.999$, $\epsilon=1e-8$):**
Achieved a **training accuracy of 96.41%**, but the **validation accuracy remained low at 12.75%**, suggesting severe overfitting.
- **High β_1 ($\beta_1=0.95$, $\beta_2=0.999$, $\epsilon=1e-8$):**
Training accuracy reached **97.20%**, with a **validation accuracy of 23.66%**, showing slightly improved generalization compared to the baseline.
- **Low β_2 ($\beta_1=0.9$, $\beta_2=0.99$, $\epsilon=1e-8$):**
Training accuracy of **95.27%** and validation accuracy of **23.66%**, similar to the high β_1 configuration.
- **Large ϵ ($\beta_1=0.9$, $\beta_2=0.999$, $\epsilon=1e-6$):**
Training accuracy of **95.99%**, but validation accuracy dropped to **11.74%**, indicating instability during optimization.
- **Very High Smoothing ($\beta_1=0.99$, $\beta_2=0.9999$, $\epsilon=1e-7$):**
Achieved training accuracy of **95.22%** and validation accuracy of **23.66%**, comparable to high β_1 and low β_2 configurations.

Observation: Across all experiments, training accuracy was high, but validation accuracy was generally low, suggesting that the model overfitted the training data. The configurations with slightly higher β_1 or lower β_2 values demonstrated better generalization.

4.2. Training Time:


The total training time varied between configurations:

Configuration	Training Time (s)
Baseline	3215.25
High β_1	2451.52
Low β_2	2442.72
Large ϵ	2463.02
Very High Smoothing	2483.80

Observation: The baseline configuration took the longest to train, possibly due to slower convergence. Other configurations achieved similar performance in shorter training times.

4.3. Graphs and Tables:

4.3.1 Validation Accuracy vs. Epochs:




Interpretation:

- The validation accuracy curves show that most configurations remained almost flat throughout all **10 epochs**.
- The baseline and large ϵ configurations stayed near **0.12**, showing very poor generalization.
- The high β_1 , low β_2 , and very high smoothing configurations performed slightly better, maintaining validation accuracy around **0.23–0.24** during the first few epochs.
- The low β_2 configuration dropped sharply after epoch 1, suggesting unstable generalization.

Overall, the plot indicates that although some hyperparameter settings improved validation accuracy slightly, all configurations demonstrated limited improvement across epochs, indicating either dataset complexity or severe overfitting.

4.3.2 Training Loss vs. Epochs



Interpretation:

- The training loss decreased steadily for all configurations, with rapid reduction in the first few epochs.
- Baseline and large ϵ converged smoothly, while very high smoothing was slightly slower initially but reached similar levels.
- By epoch 7, all loss curves were close, showing Adam reliably minimizing training loss regardless of hyperparameter variations.
- Validation accuracy varied slightly, with High β_1 , Low β_2 , and Very High Smoothing achieving the highest (23.66%), but differences were not statistically significant. Training time was shorter for all configurations except the baseline, indicating that hyperparameter tuning can improve efficiency without harming performance.
- Overall, small changes in β_1 , β_2 , and ϵ slightly affect generalization and training speed but do not significantly boost accuracy.

5.Discussion:

1. How did the different settings of Adam's hyperparameters impact model performance?

Changing the Adam hyperparameters mostly affected training behavior, not the final validation accuracy. Higher β_1 and lower β_2 made the model learn faster, while large ϵ made the training unstable and reduced performance. Overall, the default Adam settings remained the most stable.

2. Did any setting lead to faster convergence or better accuracy?

Yes the High β_1 and Low β_2 configurations converged faster than the baseline and reached high training accuracy earlier.

However, none of the configurations significantly improved validation accuracy. The highest validation accuracy still appeared in the High β_1 and Very High Smoothing settings, but the improvement was small.

3. What are the trade-offs involved with adjusting each hyperparameter?

Higher β_1 : Faster learning early on, but no major improvement in validation accuracy.

Lower β_2 : Faster convergence but weaker generalization.

Very High Smoothing ($\beta_1=0.99$, $\beta_2=0.9999$): More stable training, slightly better validation accuracy, but slower updates.

Large ϵ : Training becomes unstable and results in poor validation accuracy.

Overall, the trade-off was between speed vs. stability vs. generalization, with no setting clearly outperforming the default.

6. Conclusion:

- In conclusion, the experiments showed that Adam's hyperparameters have a noticeable impact on how the model learns, even though they did not drastically change the final validation accuracy. Increasing β_1 or lowering β_2 helped the model converge faster during training, but this improvement did not translate into better generalization on the validation set. Very high smoothing made the training more stable but did not significantly outperform the baseline. On the other hand, using a large ϵ made training unstable and led to the worst validation accuracy.
- Overall, the results highlight that adjusting Adam's hyperparameters mainly affects training speed and stability, while the standard Adam settings remain the most reliable. For future work, experiments could explore learning-rate scheduling, weight decay, or comparing Adam with other optimizers such as RMSProp or SGD with momentum.

7. References:

Below is a sample references list (modify based on what you actually used):

1. Kingma, D. P., & Ba, J. (2015). Adam: A Method for Stochastic Optimization. International Conference on Learning Representations (ICLR).
2. TensorFlow Documentation: `tf.keras.optimizers.Adam`.
3. Howard, A. et al. (2018). MobileNetV2: Inverted Residuals and Linear Bottlenecks. CVPR.
4. Chollet, F. (2017). Deep Learning with Python. Manning Publications.
5. Keras API Documentation: <https://keras.io/>

THE END